

House Prices Prediction – EDA

Performed by:

Zahra Eshtiaghi 476679

Tsoi Kwan Ma 476914

Under Supervision of Professor Paweł Sakowski

Introduction

This notebook focuses on building a regression pipeline to predict house sale prices using the **Ames Housing dataset** from the Kaggle competition *House Prices: Advanced Regression Techniques*.

The objective of this notebook is to perform **systematic data preprocessing and feature preparation** prior to model training. We begin with exploratory analysis to understand the structure of the dataset, identify missing values, and analyze feature distributions. Special attention is given to handling **skewed numerical features** and **categorical variables**, as these are common challenges in real-world tabular data.

The preprocessing pipeline includes:

- Separation of features and target variable
- Train–validation split
- Identification of numerical and categorical features
- Missing value handling
- Log transformation of skewed numerical variables
- Feature encoding

The processed datasets are then saved to disk to ensure **reproducibility and modularity**.

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import missingno as msno
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import make_pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import OneHotEncoder
```

Data Loading

In this section, we load the training and test datasets and perform an initial inspection to understand their structure. This step ensures that the data has been imported correctly and allows us to identify the number of observations and features available for modeling.

```
In [3]: df_test = pd.read_csv('data/test.csv')
df = pd.read_csv('data/train.csv') # Replace with the correct file

# Show basic info
print("\nFirst 5 rows:")
print(df.head())
```

First 5 rows:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotSha
0	1	60	RL	65.0	8450	Pave	NaN	R
1	2	20	RL	80.0	9600	Pave	NaN	R
2	3	60	RL	68.0	11250	Pave	NaN	I
3	4	70	RL	60.0	9550	Pave	NaN	I
4	5	60	RL	84.0	14260	Pave	NaN	I

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscV
0	MoSold	Lvl	AllPub	...	0	NaN	NaN	NaN
0	2	Lvl	AllPub	...	0	NaN	NaN	NaN
1	5	Lvl	AllPub	...	0	NaN	NaN	NaN
2	9	Lvl	AllPub	...	0	NaN	NaN	NaN
3	2	Lvl	AllPub	...	0	NaN	NaN	NaN
4	12	Lvl	AllPub	...	0	NaN	NaN	NaN

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

Initial Data Exploration

We explore the dataset by examining data types, summary statistics, and missing values. This analysis helps identify potential data quality issues and guides the selection of appropriate preprocessing techniques for numerical and categorical features.

In [4]: `# the structure of the dataset
df.info()`

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1460 entries, 0 to 1459  
Data columns (total 81 columns):  
 #   Column            Non-Null Count  Dtype     
---  --  
 0   Id                1460 non-null    int64    
 1   MSSubClass         1460 non-null    int64    
 2   MSZoning          1460 non-null    object    
 3   LotFrontage       1201 non-null    float64   
 4   LotArea            1460 non-null    int64    
 5   Street             1460 non-null    object    
 6   Alley              91 non-null     object    
 7   LotShape           1460 non-null    object    
 8   LandContour        1460 non-null    object    
 9   Utilities          1460 non-null    object    
 10  LotConfig          1460 non-null    object    
 11  LandSlope          1460 non-null    object    
 12  Neighborhood       1460 non-null    object    
 13  Condition1         1460 non-null    object    
 14  Condition2         1460 non-null    object    
 15  BldgType           1460 non-null    object    
 16  HouseStyle         1460 non-null    object    
 17  OverallQual        1460 non-null    int64    
 18  OverallCond        1460 non-null    int64    
 19  YearBuilt          1460 non-null    int64    
 20  YearRemodAdd       1460 non-null    int64    
 21  RoofStyle           1460 non-null    object    
 22  RoofMatl           1460 non-null    object    
 23  Exterior1st        1460 non-null    object    
 24  Exterior2nd        1460 non-null    object    
 25  MasVnrType          588 non-null     object    
 26  MasVnrArea         1452 non-null    float64   
 27  ExterQual          1460 non-null    object    
 28  ExterCond          1460 non-null    object    
 29  Foundation          1460 non-null    object    
 30  BsmtQual           1423 non-null    object    
 31  BsmtCond           1423 non-null    object    
 32  BsmtExposure        1422 non-null    object    
 33  BsmtFinType1        1423 non-null    object    
 34  BsmtFinSF1          1460 non-null    int64    
 35  BsmtFinType2        1422 non-null    object    
 36  BsmtFinSF2          1460 non-null    int64    
 37  BsmtUnfSF           1460 non-null    int64    
 38  TotalBsmtSF         1460 non-null    int64    
 39  Heating              1460 non-null    object    
 40  HeatingQC            1460 non-null    object    
 41  CentralAir          1460 non-null    object
```

```
42 Electrical    1459 non-null   object
43 1stFlrSF     1460 non-null   int64
44 2ndFlrSF     1460 non-null   int64
45 LowQualFinSF 1460 non-null   int64
46 GrLivArea    1460 non-null   int64
47 BsmtFullBath 1460 non-null   int64
48 BsmtHalfBath 1460 non-null   int64
49 FullBath     1460 non-null   int64
50 HalfBath     1460 non-null   int64
51 BedroomAbvGr 1460 non-null   int64
52 KitchenAbvGr 1460 non-null   int64
53 KitchenQual   1460 non-null   object
54 TotRmsAbvGrd 1460 non-null   int64
55 Functional    1460 non-null   object
56 Fireplaces    1460 non-null   int64
57 FireplaceQu   770 non-null   object
58 GarageType    1379 non-null   object
59 GarageYrBlt   1379 non-null   float64
60 GarageFinish   1379 non-null   object
61 GarageCars    1460 non-null   int64
62 GarageArea    1460 non-null   int64
63 GarageQual    1379 non-null   object
64 GarageCond    1379 non-null   object
65 PavedDrive    1460 non-null   object
66 WoodDeckSF   1460 non-null   int64
67 OpenPorchSF   1460 non-null   int64
68 EnclosedPorch 1460 non-null   int64
69 3SsnPorch     1460 non-null   int64
70 ScreenPorch   1460 non-null   int64
71 PoolArea      1460 non-null   int64
72 PoolQC        7 non-null    object
73 Fence          281 non-null   object
74 MiscFeature   54 non-null    object
75 MiscVal       1460 non-null   int64
76 MoSold        1460 non-null   int64
77 YrSold        1460 non-null   int64
78 SaleType       1460 non-null   object
79 SaleCondition  1460 non-null   object
80 SalePrice     1460 non-null   int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB
```

```
In [5]: # 1. Separate features and target
X = df.drop('SalePrice', axis=1)
y = df['SalePrice']
```

Train–Validation Split

To evaluate model performance reliably, the dataset is split into training and validation sets. This ensures that model evaluation is performed on unseen data and reduces the risk of overfitting during training.

```
In [6]: # 2. Split into train and test sets
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0)
```

```
X_test = df_test.copy()
```

```
In [7]: # Verify shapes  
print(f"Training set: {X_train.shape}, val set: {X_val.shape}")
```

Training set: (1168, 80), val set: (292, 80)

```
In [8]: columns_with_missings = X_train.columns[X_train.isnull().any()]  
  
# Print missing value counts  
print("Missing value counts:\n")  
print(X_train[columns_with_missings].isnull().sum())
```

Missing value counts:

```
LotFrontage      217  
Alley          1094  
MasVnrType     683  
MasVnrArea       6  
BsmtQual        28  
BsmtCond        28  
BsmtExposure    28  
BsmtFinType1    28  
BsmtFinType2    28  
Electrical        1  
FireplaceQu     547  
GarageType       64  
GarageYrBlt      64  
GarageFinish     64  
GarageQual       64  
GarageCond       64  
PoolQC          1162  
Fence           935  
MiscFeature     1122  
dtype: int64
```

```
In [9]: # Print frequency tables for each column  
print("\nFrequency counts for columns with missing values:\n")  
for column in columns_with_missings:  
    print(f"\nFrequency counts for {column}:\n")  
    print(X_train[column].value_counts(dropna=False))  
    print("-" * 50) # Add separator between columns
```

Frequency counts for columns with missing values:

Frequency counts for LotFrontage:

```
LotFrontage  
NaN      217  
60.0     112  
70.0      57  
80.0      54  
50.0      47  
...  
106.0      1  
168.0      1  
141.0      1
```

```
144.0      1  
153.0      1  
Name: count, Length: 108, dtype: int64
```

```
Frequency counts for Alley:  
Alley  
NaN      1094  
Grvl      44  
Pave      30  
Name: count, dtype: int64
```

```
Frequency counts for MasVnrType:  
MasVnrType  
NaN      683  
BrkFace    366  
Stone      106  
BrkCmn     13  
Name: count, dtype: int64
```

```
Frequency counts for MasVnrArea:  
MasVnrArea  
0.0      674  
108.0     7  
180.0     7  
NaN       6  
106.0     6  
...  
254.0     1  
860.0     1  
128.0     1  
796.0     1  
309.0     1  
Name: count, Length: 287, dtype: int64
```

```
Frequency counts for BsmtQual:  
BsmtQual  
TA      521  
Gd      493  
Ex      97  
Fa      29  
NaN     28  
Name: count, dtype: int64
```

```
Frequency counts for BsmtCond:  
BsmtCond  
TA      1045  
Gd      55  
Fa      39  
NaN     28  
Po      1  
Name: count, dtype: int64
```

Frequency counts for BsmtExposure:

BsmtExposure

No	769
Av	175
Gd	103
Mn	93
NaN	28

Name: count, dtype: int64

Frequency counts for BsmtFinType1:

BsmtFinType1

Unf	345
GLQ	328
ALQ	178
BLQ	123
Rec	104
LwQ	62
NaN	28

Name: count, dtype: int64

Frequency counts for BsmtFinType2:

BsmtFinType2

Unf	1009
Rec	46
LwQ	36
NaN	28
BLQ	22
ALQ	17
GLQ	10

Name: count, dtype: int64

Frequency counts for Electrical:

Electrical

SBrkr	1071
FuseA	69
FuseF	24
FuseP	3
NaN	1

Name: count, dtype: int64

Frequency counts for FireplaceQu:

FireplaceQu

NaN	547
Gd	305
TA	252
Fa	27
Ex	21
Po	16

Name: count, dtype: int64

Frequency counts for GarageType:

GarageType

Attchd	693
Detchd	308
BuiltIn	74
Nan	64
Basment	16
CarPort	7
2Types	6

Name: count, dtype: int64

Frequency counts for GarageYrBlt:

GarageYrBlt

Nan	64
2005.0	52
2004.0	45
2006.0	45
2003.0	40
..	
1922.0	2
1934.0	2
1908.0	1
1900.0	1
1906.0	1

Name: count, Length: 95, dtype: int64

Frequency counts for GarageFinish:

GarageFinish

Unf	480
RFn	339
Fin	285
Nan	64

Name: count, dtype: int64

Frequency counts for GarageQual:

GarageQual

TA	1050
Nan	64
Fa	36
Gd	13
Ex	3
Po	2

Name: count, dtype: int64

Frequency counts for GarageCond:

GarageCond

TA	1061
Nan	64
Fa	27
Gd	9
Po	5

```
Ex      2  
Name: count, dtype: int64
```

```
Frequency counts for PoolQC:  
PoolQC  
NaN    1162  
Fa      2  
Ex      2  
Gd      2  
Name: count, dtype: int64
```

```
Frequency counts for Fence:  
Fence  
NaN    935  
MnPrv   128  
GdPrv   50  
GdWo    46  
MnWw    9  
Name: count, dtype: int64
```

```
Frequency counts for MiscFeature:  
MiscFeature  
NaN    1122  
Shed    41  
0thr    2  
Gar2    2  
TenC    1  
Name: count, dtype: int64
```

Missing Value Handling Summary:

1. **Alley, MasVnrType, BsmtQual, etc. (13 columns)** - `NaN` means feature doesn't exist → Fill with "None" (categorical placeholder).
2. **GarageYrBlt** - Missing garage year → Fill with most common year from training data (mode).
3. **MasVnrArea** - No masonry veneer → Fill with 0 square feet.
4. **LotFrontage** - Unknown street frontage → Fill with median of same neighborhood (houses nearby have similar lots).
5. **Electrical** - One missing electrical system → Fill with most common type (mode).
6. **PoolQC** - No pool → Fill with "No" (adds as new quality level).
7. **MiscFeature** - No miscellaneous feature → Fill with "None".

Logic: Preserve information where `NaN` has meaning ("doesn't exist"), use data patterns (neighborhood medians) where sensible, and fall back to common values (mode/median) for true unknowns.

```
In [10]: # For columns where NaN means "None" (categorical)
none_cols = ['Alley', 'MasVnrType', 'BsmtQual', 'BsmtCond', 'BsmtEx
            'BsmtFinType1', 'BsmtFinType2', 'FireplaceQu', 'Garage
            'GarageFinish', 'GarageQual', 'GarageCond', 'Fence']

X_train[none_cols] = X_train[none_cols].fillna('None')
X_val[none_cols] = X_val[none_cols].fillna("None")
X_test[none_cols] = X_test[none_cols].fillna('None')
```

```
In [11]: garage_mode = X_train['GarageYrBlt'].mode()[0]
X_train['GarageYrBlt'] = X_train['GarageYrBlt'].fillna(garage_mode)
X_val['GarageYrBlt'] = X_val['GarageYrBlt'].fillna(garage_mode)
X_test['GarageYrBlt'] = X_test['GarageYrBlt'].fillna(garage_mode)
```

```
In [12]: # For MasVnrArea: NaN likely means no veneer. Fill with 0

X_train['MasVnrArea'] = X_train['MasVnrArea'].fillna(0)
X_val['MasVnrArea'] = X_val['MasVnrArea'].fillna(0)
X_test['MasVnrArea'] = X_test['MasVnrArea'].fillna(0)
```

This imputes missing **LotFrontage** values with the median LotFrontage of the same neighborhood.

Why by neighborhood? Houses in the same neighborhood tend to have similar lot characteristics. This is more accurate than using the overall median

```
In [13]: # LotFrontage (median by neighborhood from TRAINING only)
lot_median_by_neigh = X_train.groupby('Neighborhood')['LotFrontage']
X_train['LotFrontage'] = X_train.apply(
    lambda row: lot_median_by_neigh[row['Neighborhood']] if pd.isna
    axis=1
)
X_val['LotFrontage'] = X_val.apply(
    lambda row: lot_median_by_neigh.get(row['Neighborhood'], lot_me
    axis=1
)
X_test['LotFrontage'] = X_test.apply(
    lambda row: lot_median_by_neigh.get(row['Neighborhood'], lot_me
    axis=1
)
```

```
In [14]: # Electrical (mode)
elec_mode = X_train['Electrical'].mode()[0]
X_train['Electrical'] = X_train['Electrical'].fillna(elec_mode)
X_val['Electrical'] = X_val['Electrical'].fillna(elec_mode)
X_test['Electrical'] = X_test['Electrical'].fillna(elec_mode)
```

```
In [15]: # PoolQC: NaN = No Pool
```

```
X_train['PoolQC'] = X_train['PoolQC'].fillna('No')
X_val['PoolQC'] = X_val['PoolQC'].fillna('No')
X_test['PoolQC'] = X_test['PoolQC'].fillna('No')
```

In [16]: # MiscFeature: NaN = None

```
X_train['MiscFeature'] = X_train['MiscFeature'].fillna('None')
X_val['MiscFeature'] = X_val['MiscFeature'].fillna('None')
X_test['MiscFeature'] = X_test['MiscFeature'].fillna('None')
```

In [17]: test_columns_with_missings = X_test.columns[X_test.isnull().any()]

```
# Print missing value counts
print("Missing value counts:\n")
print(X_test[test_columns_with_missings].isnull().sum())
```

Missing value counts:

LotFrontage	0
Alley	0
MasVnrType	0
MasVnrArea	0
BsmtQual	0
BsmtCond	0
BsmtExposure	0
BsmtFinType1	0
BsmtFinType2	0
Electrical	0
FireplaceQu	0
GarageType	0
GarageYrBlt	0
GarageFinish	0
GarageQual	0
GarageCond	0
PoolQC	0
Fence	0
MiscFeature	0

dtype: int64

In [18]: # Print frequency tables for each column

```
print("\nFrequency counts for columns with missing values after imputation")
for column in columns_with_missings:
    print(f"\nFrequency counts for {column}:")
    print(X_train[column].value_counts(dropna=False))
    print("-" * 50) # Add separator between columns
```

Frequency counts for columns with missing values after imputation:

Frequency counts for LotFrontage:

LotFrontage	
60.0	120
80.0	91
70.0	76
65.0	60
50.0	47
...	

```
144.0      1  
128.0      1  
39.0       1  
141.0      1  
153.0      1  
Name: count, Length: 114, dtype: int64
```

```
Frequency counts for Alley:  
Alley  
None    1094  
Grvl     44  
Pave     30  
Name: count, dtype: int64
```

```
Frequency counts for MasVnrType:  
MasVnrType  
None    683  
BrkFace   366  
Stone     106  
BrkCmn     13  
Name: count, dtype: int64
```

```
Frequency counts for MasVnrArea:  
MasVnrArea  
0.0      680  
108.0     7  
180.0     7  
200.0     6  
106.0     6  
...  
366.0     1  
576.0     1  
768.0     1  
378.0     1  
309.0     1  
Name: count, Length: 286, dtype: int64
```

```
Frequency counts for BsmtQual:  
BsmtQual  
TA      521  
Gd      493  
Ex      97  
Fa      29  
None    28  
Name: count, dtype: int64
```

```
Frequency counts for BsmtCond:  
BsmtCond  
TA      1045  
Gd      55  
Fa      39
```

```
None      28
Po        1
Name: count, dtype: int64
```

```
-----
```

Frequency counts for BsmtExposure:

BsmtExposure

```
No      769
Av      175
Gd      103
Mn      93
None    28
Name: count, dtype: int64
```

```
-----
```

Frequency counts for BsmtFinType1:

BsmtFinType1

```
Unf     345
GLQ     328
ALQ     178
BLQ     123
Rec     104
LwQ     62
None    28
Name: count, dtype: int64
```

```
-----
```

Frequency counts for BsmtFinType2:

BsmtFinType2

```
Unf     1009
Rec     46
LwQ     36
None    28
BLQ     22
ALQ     17
GLQ     10
Name: count, dtype: int64
```

```
-----
```

Frequency counts for Electrical:

Electrical

```
SBrkr   1072
FuseA    69
FuseF    24
FuseP    3
Name: count, dtype: int64
```

```
-----
```

Frequency counts for FireplaceQu:

FireplaceQu

```
None    547
Gd     305
TA     252
Fa     27
Ex     21
Po     16
```

Name: count, dtype: int64

Frequency counts for GarageType:

GarageType

Attchd	693
Detchd	308
BuiltIn	74
None	64
Basment	16
CarPort	7
2Types	6

Name: count, dtype: int64

Frequency counts for GarageYrBlt:

GarageYrBlt

2005.0	116
2006.0	45
2004.0	45
2003.0	40
2007.0	40
...	
1914.0	2
1934.0	2
1908.0	1
1900.0	1
1906.0	1

Name: count, Length: 94, dtype: int64

Frequency counts for GarageFinish:

GarageFinish

Unf	480
RFn	339
Fin	285
None	64

Name: count, dtype: int64

Frequency counts for GarageQual:

GarageQual

TA	1050
None	64
Fa	36
Gd	13
Ex	3
Po	2

Name: count, dtype: int64

Frequency counts for GarageCond:

GarageCond

TA	1061
None	64
Fa	27

```
Gd      9  
Po      5  
Ex      2  
Name: count, dtype: int64
```

```
-----  
Frequency counts for PoolQC:  
PoolQC  
No    1162  
Fa     2  
Ex     2  
Gd     2  
Name: count, dtype: int64
```

```
-----  
Frequency counts for Fence:  
Fence  
None    935  
MnPrv   128  
GdPrv   50  
GdWo    46  
MnWw    9  
Name: count, dtype: int64
```

```
-----  
Frequency counts for MiscFeature:  
MiscFeature  
None    1122  
Shed    41  
Othr    2  
Gar2    2  
TenC    1  
Name: count, dtype: int64
```

```
In [19]: # Verify imputation by checking for remaining NaN values  
print("NaN counts after imputation:")  
# Verify  
print("Missing in X_train:", X_train.isnull().sum().sum())  
print("Missing in X_val:", X_val.isnull().sum().sum())  
print("Missing in X_test:", X_test.isnull().sum().sum())
```

```
Nan counts after imputation:  
Missing in X_train: 0  
Missing in X_val: 0  
Missing in X_test: 22
```

```
In [20]: test_columns_with_missings = X_test.columns[X_test.isnull().any()]  
  
# Print missing value counts  
print("Missing value counts:\n")  
print(X_test[test_columns_with_missings].isnull().sum())
```

Missing value counts:

```
MSZoning      4
Utilities     2
Exterior1st   1
Exterior2nd   1
BsmtFinSF1    1
BsmtFinSF2    1
BsmtUnfSF     1
TotalBsmtSF   1
BsmtFullBath  2
BsmtHalfBath  2
KitchenQual   1
Functional    2
GarageCars    1
GarageArea    1
SaleType       1
dtype: int64
```

The test set cannot have missing values for prediction.

- For categorical: Use training mode
- For basement/garage numerical: If corresponding quality is "None", fill with 0
- This maintains data consistency

```
In [21]: # Categorical columns - fill with mode from TRAINING
cat_cols = ['MSZoning', 'Utilities', 'Exterior1st', 'Exterior2nd',
            'KitchenQual', 'Functional', 'SaleType']

for col in cat_cols:
    mode_val = X_train[col].mode()[0]
    X_test[col] = X_test[col].fillna(mode_val)

# Numerical basement columns - if basement features missing, likely
# Check if corresponding BsmtQual is 'None'
bsmt_num_cols = ['BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF',
                  'BsmtFullBath', 'BsmtHalfBath']

for col in bsmt_num_cols:
    # If basement quality is 'None', fill with 0
    mask = (X_test['BsmtQual'] == 'None') & (X_test[col].isna())
    X_test.loc[mask, col] = 0

    # For any remaining NaN (shouldn't happen), fill with 0
    X_test[col] = X_test[col].fillna(0)

# Garage columns - similar logic
garage_num_cols = ['GarageCars', 'GarageArea']

for col in garage_num_cols:
    # If garage type is 'None', fill with 0
    mask = (X_test['GarageType'] == 'None') & (X_test[col].isna())
    X_test.loc[mask, col] = 0
```

```
# For any remaining NaN, fill with 0
X_test[col] = X_test[col].fillna(0)

# Verify no missing values remain
print("Remaining missing in X_test:", X_test.isnull().sum().sum())
```

Remaining missing in X_test: 0

Feature Type Identification

Features are divided into **numerical** and **categorical** groups. This separation allows us to apply appropriate preprocessing techniques to each feature type, such as imputation and scaling for numerical variables and encoding for categorical variables.

In [22]: # Lets check the frequencies of individual levels of all the nominal variables

```
# Based on the information about the dataset we can list all nominal variables =
['MSSubClass',      # Dwelling type – categories not ordered
 'MSZoning',        # Zoning classification
 'Street',          # Road type
 'Alley',           # Alley type
 'LotShape',         # Shape categories not ordered
 'LandContour',     # Flatness types
 'Utilities',        # Utility types
 'LotConfig',        # Lot configuration
 'LandSlope',        # Slope categories
 'Neighborhood',    # Location names
 'Condition1',      # Proximity conditions
 'Condition2',      # Secondary conditions
 'BldgType',         # Dwelling type
 'HouseStyle',       # Style categories not ordered
 'RoofStyle',        # Roof types
 'RoofMatl',         # Roof materials
 'Exterior1st',     # Exterior covering
 'Exterior2nd',     # Secondary exterior
 'MasVnrType',       # Masonry type
 'Foundation',       # Foundation type
 'Heating',          # Heating type
 'Electrical',       # Electrical system
 'Functional',       # Functionality categories
 'GarageType',       # Garage location types
 'GarageFinish',     # Finish categories (could be ordinal but treated as nominal)
 'PavedDrive',       # Driveway type
 'MiscFeature',      # Miscellaneous features
 'SaleType',          # Sale type
 'SaleCondition'     # Sale condition
]
```

In [23]: # Loop through each nominal variable

```
for var in nominal_variables:
    print(f"\n{var}:")
    print(X_train[var].value_counts(dropna=False))
```

```
    print("-" * 50)

MSSubClass:
MSSubClass
20      434
60      240
50      113
120     64
70      52
30      50
160     49
80      45
90      41
190     28
85      17
75      15
45      10
180     7
40      3
Name: count, dtype: int64
```

```
MSZoning:
MSZoning
RL      924
RM      172
FV      53
RH      15
C (all)  4
Name: count, dtype: int64
```

```
Street:
Street
Pave    1164
Grvl     4
Name: count, dtype: int64
```

```
Alley:
Alley
None    1094
Grvl     44
Pave     30
Name: count, dtype: int64
```

```
LotShape:
LotShape
Reg     729
IR1    394
IR2     37
IR3     8
Name: count, dtype: int64
```

LandContour:
LandContour
Lvl 1059
Bnk 48
HLS 35
Low 26
Name: count, dtype: int64

Utilities:
Utilities
AllPub 1167
NoSeWa 1
Name: count, dtype: int64

LotConfig:
LotConfig
Inside 822
Corner 221
CulDSac 84
FR2 38
FR3 3
Name: count, dtype: int64

LandSlope:
LandSlope
Gtl 1108
Mod 51
Sev 9
Name: count, dtype: int64

Neighborhood:
Neighborhood
NAmes 181
CollgCr 115
OldTown 91
Edwards 87
Somerst 69
NWAmes 66
Gilbert 65
NridgHt 61
Sawyer 58
BrkSide 45
Crawfor 44
SawyerW 44
Mitchel 40
NoRidge 33
Timber 28
IDOTRR 26
SWISU 21
StoneBr 20
ClearCr 19
Blmngtn 15

```
BrDale      13
MeadowV     10
Veenker      9
NPKVill      7
Blueste      1
Name: count, dtype: int64
```

```
Condition1:
Condition1
Norm      1004
Feedr     66
Artery    40
RRAn      19
PosN      15
RRAe      10
PosA      8
RRNn      5
RRNe      1
Name: count, dtype: int64
```

```
Condition2:
Condition2
Norm      1157
Feedr     3
PosN      2
Artery    2
RRNn      1
PosA      1
RRAe      1
RRAn      1
Name: count, dtype: int64
```

```
BldgType:
BldgType
1Fam      978
TwnhsE    88
Duplex    41
Twnhs     32
2fmCon    29
Name: count, dtype: int64
```

```
HouseStyle:
HouseStyle
1Story    577
2Story    360
1.5Fin    121
SLvl      52
SFoyer    28
1.5Unf    12
2.5Unf    11
2.5Fin    7
Name: count, dtype: int64
```

```
RoofStyle:  
RoofStyle  
Gable      906  
Hip        235  
Flat       11  
Gambrel     9  
Mansard     5  
Shed        2  
Name: count, dtype: int64
```

```
RoofMatl:  
RoofMatl  
CompShg    1149  
Tar&Grv      9  
WdShngl     4  
WdShake     3  
Metal       1  
ClyTile     1  
Roll        1  
Name: count, dtype: int64
```

```
Exterior1st:  
Exterior1st  
VinylSd    420  
HdBoard    176  
MetalSd    173  
Wd Sdng    171  
Plywood     81  
CemntBd    45  
BrkFace     40  
Stucco      21  
WdShing     19  
AsbShng     16  
BrkComm      2  
ImStucc      1  
CBlock       1  
AsphShn      1  
Stone        1  
Name: count, dtype: int64
```

```
Exterior2nd:  
Exterior2nd  
VinylSd    410  
MetalSd    165  
Wd Sdng    165  
HdBoard    163  
Plywood     112  
CmentBd    44  
Wd Shng     34  
Stucco      23  
AsbShng     17
```

```
BrkFace      15
ImStucc       6
Brk Cmn       6
AsphShn        3
Stone          3
Other          1
CBlock         1
Name: count, dtype: int64
```

```
MasVnrType:
MasVnrType
None      683
BrkFace    366
Stone      106
BrkCmn     13
Name: count, dtype: int64
```

```
Foundation:
Foundation
PConc      520
CBlock     504
BrkTil     116
Slab        20
Stone        5
Wood        3
Name: count, dtype: int64
```

```
Heating:
Heating
GasA      1140
GasW       15
Grav        6
Wall        4
OthW        2
Floor       1
Name: count, dtype: int64
```

```
Electrical:
Electrical
SBrkr     1072
FuseA      69
FuseF      24
FuseP      3
Name: count, dtype: int64
```

```
Functional:
Functional
Typ      1084
Min2      29
Min1      28
Mod       13
```

```
Maj1      9  
Maj2      4  
Sev       1  
Name: count, dtype: int64
```

```
GarageType:  
GarageType  
Attchd    693  
Detchd    308  
BuiltIn    74  
None      64  
Basment   16  
CarPort    7  
2Types     6  
Name: count, dtype: int64
```

```
GarageFinish:  
GarageFinish  
Unf      480  
RFn      339  
Fin      285  
None     64  
Name: count, dtype: int64
```

```
PavedDrive:  
PavedDrive  
Y      1070  
N      73  
P      25  
Name: count, dtype: int64
```

```
MiscFeature:  
MiscFeature  
None    1122  
Shed     41  
Othr     2  
Gar2     2  
TenC     1  
Name: count, dtype: int64
```

```
SaleType:  
SaleType  
WD      1012  
New     97  
COD     36  
ConLD    7  
ConLI    4  
ConLw    4  
CWD     4  
Oth     2  
Con     2
```

```
Name: count, dtype: int64
```

```
SaleCondition:
```

```
SaleCondition
```

```
Normal    964
```

```
Partial    98
```

```
Abnorml    77
```

```
Family     18
```

```
Alloca      7
```

```
AdjLand     4
```

```
Name: count, dtype: int64
```

```
In [24]: # Threshold: categories with less than 1% of data (approx 11 houses
threshold = 0.01 * len(X_train)

# Columns to NOT group (keep all categories)
do_not_group = ['Street', 'Utilities', 'Condition2'] # Add others

for var in nominal_variables:
    if var in do_not_group:
        continue # Skip grouping for these

    value_counts = X_train[var].value_counts()

    if 'None' in value_counts.index:
        continue

    rare_categories = value_counts[value_counts < threshold].index

    if len(rare_categories) > 0:
        X_train[var] = X_train[var].replace(rare_categories, 'Other')
        X_val[var] = X_val[var].replace(rare_categories, 'Other')
        X_test[var] = X_test[var].replace(rare_categories, 'Other')

    print(f"{var}: Grouped {len(rare_categories)} rare categories")
```

```
MSSubClass: Grouped 3 rare categories
MSZoning: Grouped 1 rare categories
LotShape: Grouped 1 rare categories
LandContour: Grouped 0 rare categories
LotConfig: Grouped 1 rare categories
LandSlope: Grouped 1 rare categories
Neighborhood: Grouped 4 rare categories
Condition1: Grouped 4 rare categories
BldgType: Grouped 0 rare categories
HouseStyle: Grouped 2 rare categories
RoofStyle: Grouped 4 rare categories
RoofMatl: Grouped 6 rare categories
Exterior1st: Grouped 5 rare categories
Exterior2nd: Grouped 6 rare categories
Foundation: Grouped 2 rare categories
Heating: Grouped 4 rare categories
Electrical: Grouped 1 rare categories
Functional: Grouped 3 rare categories
PavedDrive: Grouped 0 rare categories
SaleType: Grouped 6 rare categories
SaleCondition: Grouped 2 rare categories
```

```
In [25]: # Check what was grouped for key columns
important_cols = ['MSZoning', 'Neighborhood', 'Exterior1st', 'Exterior2nd', 'Foundation', 'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'PavedDrive', 'SaleType', 'SaleCondition']

for var in important_cols:
    print(f"\n{var} categories after grouping:")
    print(X_train[var].value_counts())
    print("-" * 50)
```

```
MSZoning categories after grouping:
```

```
MSZoning
RL      924
RM      172
FV       53
RH       15
Other      4
Name: count, dtype: int64
```

```
Neighborhood categories after grouping:
```

```
Neighborhood
NAmes     181
CollgCr   115
OldTown    91
Edwards    87
Somerst    69
NWAmes     66
Gilbert    65
NridgHt    61
Sawyer     58
BrkSide    45
Crawfor    44
SawyerW    44
Mitchel    40
NoRidge    33
```

```
Timber      28
Other       27
IDOTRR     26
SWISU       21
StoneBr    20
ClearCr    19
Blmngtn   15
BrDale     13
Name: count, dtype: int64
```

Exterior1st categories after grouping:

```
Exterior1st
VinylSd    420
HdBoard   176
MetalSd   173
Wd Sdng   171
Plywood    81
CemntBd   45
BrkFace    40
Stucco     21
WdShing   19
AsbShng   16
Other      6
Name: count, dtype: int64
```

Exterior2nd categories after grouping:

```
Exterior2nd
VinylSd    410
MetalSd   165
Wd Sdng   165
HdBoard   163
Plywood    112
CmentBd   44
Wd Shng   34
Stucco     23
Other      20
AsbShng   17
BrkFace    15
Name: count, dtype: int64
```

SaleType categories after grouping:

```
SaleType
WD        1012
New       97
COD       36
Other     23
Name: count, dtype: int64
```

The test/val sets might have new categories not seen in training.

lets Check for unseen categories:

```
In [26]: # After grouping, check for categories in test/val that don't exist
for var in nominal_variables:
    train_cats = set(X_train[var].unique())
    val_cats = set(X_val[var].unique())
    test_cats = set(X_test[var].unique())

    # Find categories in val/test but not in training
    val_unseen = val_cats - train_cats
    test_unseen = test_cats - train_cats

    if val_unseen:
        print(f"\n{var}: Unseen in val - {val_unseen}")
        # Replace unseen with 'Other' or most common
        X_val.loc[X_val[var].isin(val_unseen), var] = 'Other'

    if test_unseen:
        print(f"\n{var}: Unseen in test - {test_unseen}")
        X_test.loc[X_test[var].isin(test_unseen), var] = 'Other'
```

MSSubClass: Unseen in test - {150}
RoofMatl: Unseen in val - {'Membran'}
Electrical: Unseen in val - {'Mix'}

These are new categories in test/validation that don't exist in training:

```
In [27]: # For MSSubClass: 150 is "1-1/2 STORY PUD - ALL AGES"
# Since it's unseen, map it to the most similar category or 'Other'
# Check what's most common PUD type in training
pud_classes = [120, 150, 160, 180] # All PUD types
common_pud = X_train[X_train['MSSubClass'].isin(pud_classes)]['MSSubClass']
if len(common_pud) > 0:
    replacement = common_pud[0] # Most common PUD type
else:
    replacement = 'Other' # Fallback

X_test.loc[X_test['MSSubClass'] == 150, 'MSSubClass'] = replacement

# For RoofMatl: 'Membran' is membrane roofing
# Since extremely rare, map to 'Other' (already grouped other rare
X_val.loc[X_val['RoofMatl'] == 'Membran', 'RoofMatl'] = 'Other'

# For Electrical: 'Mix' is mixed system
# Map to most common electrical type
elec_mode = X_train['Electrical'].mode()[0]
X_val.loc[X_val['Electrical'] == 'Mix', 'Electrical'] = elec_mode
```

```
In [28]: # Verify all categories now match
for var in ['MSSubClass', 'RoofMatl', 'Electrical']:
    train_cats = set(X_train[var].unique())
    val_cats = set(X_val[var].unique())
    test_cats = set(X_test[var].unique())

    print(f"\n{var}:")
    print(f"Train: {len(train_cats)} categories")
    print(f"Val: {len(val_cats)} categories, matches: {val_cats.issubset(train_cats)}")
    print(f"Test: {len(test_cats)} categories, matches: {test_cats.issubset(train_cats)}")
```

```
MSSubClass:  
Train: 13 categories  
Val: 13 categories, matches: True  
Test: 13 categories, matches: True  
  
RoofMatl:  
Train: 2 categories  
Val: 2 categories, matches: True  
Test: 2 categories, matches: True  
  
Electrical:  
Train: 4 categories  
Val: 4 categories, matches: True  
Test: 4 categories, matches: True
```

Near-zero variance

```
In [29]: # lets define a function checking near-zero variance  
def near_zero_var(df, freq_cut=95/5, unique_cut=10):  
    """  
        Identifies columns with near-zero variance in a DataFrame and c  
  
    Parameters:  
    - df (pd.DataFrame): Input DataFrame.  
    - freq_cut (float): Threshold for the frequency ratio (default :  
    - unique_cut (int): Threshold for the unique value ratio (defau  
  
    Returns:  
    - pd.DataFrame: A sorted DataFrame containing:  
        - variable: Column name  
        - freq_ratio: Ratio of the most common value to the second  
        - unique_ratio: Ratio of unique values to total observation  
        - high_freq_ratio: Binary indicator (1 if freq_ratio > freq  
        - low_unique_ratio: Binary indicator (1 if unique_ratio < u  
    """  
    results = []  
  
    for col in df.columns:  
        # Get the value counts  
        counts = df[col].value_counts()  
  
        # Calculate freq_ratio  
        if len(counts) > 1:  
            freq_ratio = counts.iloc[0] / counts.iloc[1]  
        else:  
            freq_ratio = float('inf') # Only one unique value  
  
        # Calculate unique_ratio  
        unique_ratio = len(counts) / len(df)  
  
        # Determine binary indicators  
        high_freq_ratio = int(freq_ratio > freq_cut)  
        low_unique_ratio = int(unique_ratio < unique_cut)  
        results.append([col, freq_ratio, unique_ratio, high_freq_ratio, low_unique_ratio])  
    return pd.DataFrame(results, columns=['variable', 'freq_ratio', 'unique_ratio', 'high_freq_ratio', 'low_unique_ratio'])
```

```

# Append results
results.append({
    'variable': col,
    'freq_ratio': freq_ratio,
    'unique_ratio': unique_ratio,
    'high_freq_ratio': high_freq_ratio,
    'low_unique_ratio': low_unique_ratio
})

# Convert results to a DataFrame
results_df = pd.DataFrame(results)

# Sort by 'high_freq_ratio' (descending) and 'low_unique_ratio'
results_df = results_df.sort_values(by=['freq_ratio', 'unique_r

return results_df

```

In [30]: # Apply to X_train (training features after preprocessing)
nzv_df = near_zero_var(X_train, freq_cut=95/5, unique_cut=10)

```

# Get near-zero variance columns
nzv_columns = nzv_df[(nzv_df['low_unique_ratio'] == 1) &
                     (nzv_df['high_freq_ratio'] == 1)]['variable'].

print(f"Near-zero variance columns: {nzv_columns}")
print(f"Count: {len(nzv_columns)}")

```

Near-zero variance columns: ['Utilities', 'PoolArea', 'PoolQC', '3Ss
nPorch', 'Condition2', 'LowQualFinSF', 'Street', 'BsmtFinSF2', 'Scre
enPorch', 'MiscVal', 'MasVnrArea', '2ndFlrSF', 'EnclosedPorch', 'Hea
ting', 'RoofMatl', 'Functional', 'BsmtFinSF1', 'MiscFeature', 'Alle
y', 'LandContour', 'BsmtFinType2', 'LandSlope', 'OpenPorchSF', 'Wood
DeckSF', 'KitchenAbvGr']
Count: 25

In [31]: # For each nzv column, check if it has ANY predictive power
nzv_to_keep = []
for col in nzv_columns:
 if col in X_train.columns: # Ensure column exists
 # Simple check: mean target by category
 if X_train[col].nunique() > 1: # Has variation
 temp_df = pd.DataFrame({'feature': X_train[col], 'target': y_train})
 group_means = temp_df.groupby('feature')['target'].mean()
 # If categories have different mean prices, keep
 if group_means.std() > y_train.std() * 0.05: # 5% of o
 nzv_to_keep.append(col)
 print(f"Keep {col}: price variation = {group_means.std()}")

Drop only truly useless columns
cols_to_drop = [col for col in nzv_columns if col not in nzv_to_ke
print(f"\nDropping: {cols_to_drop}")

X_train = X_train.drop(columns=cols_to_drop)
X_val = X_val.drop(columns=cols_to_drop, errors='ignore')
X_test = X_test.drop(columns=cols_to_drop, errors='ignore')

```

Keep Utilities: price variation = 31098
Keep PoolArea: price variation = 205176
Keep PoolQC: price variation = 143696
Keep 3SsnPorch: price variation = 77937
Keep Condition2: price variation = 80240
Keep LowQualFinSF: price variation = 83969
Keep Street: price variation = 14477
Keep BsmtFinSF2: price variation = 53812
Keep ScreenPorch: price variation = 77219
Keep MiscVal: price variation = 51678
Keep MasVnrArea: price variation = 83432
Keep 2ndFlrSF: price variation = 86202
Keep EnclosedPorch: price variation = 53678
Keep Heating: price variation = 50462
Keep RoofMatl: price variation = 28206
Keep Functional: price variation = 19791
Keep BsmtFinSF1: price variation = 81784
Keep MiscFeature: price variation = 55742
Keep Alley: price variation = 31040
Keep LandContour: price variation = 42719
Keep BsmtFinType2: price variation = 34148
Keep LandSlope: price variation = 22690
Keep OpenPorchSF: price variation = 65857
Keep WoodDeckSF: price variation = 84806
Keep KitchenAbvGr: price variation = 32083

```

Dropping: []

Summary: Some features look “rare” or low-variance (e.g., `PoolQC`, `Street`, `Utilities`, `Condition2`), but they show **large price differences** across categories. This means they carry **strong predictive signal** despite low frequency.

Key Point: Low variance ≠ low importance. Even uncommon categories can influence price significantly (e.g., houses with pools are much more expensive; gravel streets are cheaper).

Conclusion: I should not drop these features. They are rare but informative.

Next Step – Encoding:

- **Binary features:** simple 0/1 encoding
- **Numeric counts:** keep as continuous values
- **Ordinal categories:** label encode
- **Nominal categories:** one-hot encode (rare levels → “Other”)

One Hot Encoding

```
In [32]: # 1. BINARY FEATURES (2 categories)
binary_cols = []
for col in X_train.columns:
    if X_train[col].nunique() == 2:
```

```

    binary_cols.append(col)

    print(f"Binary columns: {binary_cols}")

Binary columns: ['Street', 'Utilities', 'RoofMatl', 'CentralAir']

```

```

In [33]: # Encode binary (0/1)
for col in binary_cols:
    # Get the two categories
    categories = X_train[col].unique()
    if len(categories) == 2:
        # Map to 0 and 1 (alphabetical order)
        mapping = {categories[0]: 0, categories[1]: 1}
        X_train[col] = X_train[col].map(mapping)
        X_val[col] = X_val[col].map(mapping)
        X_test[col] = X_test[col].map(mapping)

```

```

In [34]: # 2. ORDINAL CATEGORICAL (inherent order)
ordinal_mappings = {
    'ExterQual': {'Po': 1, 'Fa': 2, 'TA': 3, 'Gd': 4, 'Ex': 5},
    'ExterCond': {'Po': 1, 'Fa': 2, 'TA': 3, 'Gd': 4, 'Ex': 5},
    'BsmtQual': {'None': 0, 'Po': 1, 'Fa': 2, 'TA': 3, 'Gd': 4, 'Ex': 5},
    'BsmtCond': {'None': 0, 'Po': 1, 'Fa': 2, 'TA': 3, 'Gd': 4, 'Ex': 5},
    'BsmtExposure': {'None': 0, 'No': 1, 'Mn': 2, 'Av': 3, 'Gd': 4},
    'BsmtFinType1': {'None': 0, 'Unf': 1, 'LwQ': 2, 'Rec': 3, 'BLQ': 4},
    'BsmtFinType2': {'None': 0, 'Unf': 1, 'LwQ': 2, 'Rec': 3, 'BLQ': 4},
    'HeatingQC': {'Po': 1, 'Fa': 2, 'TA': 3, 'Gd': 4, 'Ex': 5},
    'KitchenQual': {'Po': 1, 'Fa': 2, 'TA': 3, 'Gd': 4, 'Ex': 5},
    'FireplaceQu': {'None': 0, 'Po': 1, 'Fa': 2, 'TA': 3, 'Gd': 4},
    'GarageFinish': {'None': 0, 'Unf': 1, 'RFn': 2, 'Fin': 3},
    'GarageQual': {'None': 0, 'Po': 1, 'Fa': 2, 'TA': 3, 'Gd': 4, 'Ex': 5},
    'GarageCond': {'None': 0, 'Po': 1, 'Fa': 2, 'TA': 3, 'Gd': 4, 'Ex': 5},
    'PoolQC': {'No': 0, 'Fa': 1, 'TA': 2, 'Gd': 3, 'Ex': 4},
    'Fence': {'None': 0, 'MnWw': 1, 'GdWo': 2, 'MnPrv': 3, 'GdPrv': 4},
    'Functional': {'Sal': 1, 'Sev': 2, 'Maj2': 3, 'Maj1': 4, 'Mod': 5}
}

```

```

In [35]: # Apply ordinal encoding
for col, mapping in ordinal_mappings.items():
    if col in X_train.columns:
        X_train[col] = X_train[col].map(mapping)
        X_val[col] = X_val[col].map(mapping)
        X_test[col] = X_test[col].map(mapping)

```

```

In [36]: # Explicit nominal columns (based on your dataset description)
explicit_nominal = [
    'MSSubClass', 'MSZoning', 'Alley', 'LotShape', 'LandContour',
    'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'BldgType',
    'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd',
    'MasVnrType', 'Foundation', 'Heating', 'Electrical', 'GarageType',
    'PavedDrive', 'MiscFeature', 'SaleType', 'SaleCondition'
]

# Filter to only those that exist in data
nominal_cols = [col for col in explicit_nominal if col in X_train.c

```

```
print(f"One-hot encoding {len(nominal_cols)} nominal columns")
```

One-hot encoding 24 nominal columns

```
In [37]: # NOMINAL CATEGORICAL (one-hot encode)
# Convert all nominal columns to string type
for col in nominal_cols:
    X_train[col] = X_train[col].astype(str)
    X_val[col] = X_val[col].astype(str)
    X_test[col] = X_test[col].astype(str)

# Now one-hot encode
from sklearn.preprocessing import OneHotEncoder

encoder = OneHotEncoder(sparse_output=False, handle_unknown='ignore')
encoder.fit(X_train[nominal_cols])

# Transform
train_encoded = encoder.transform(X_train[nominal_cols])
val_encoded = encoder.transform(X_val[nominal_cols])
test_encoded = encoder.transform(X_test[nominal_cols])

# Create DataFrames
encoded_cols = encoder.get_feature_names_out(nominal_cols)
train_encoded_df = pd.DataFrame(train_encoded, columns=encoded_cols)
val_encoded_df = pd.DataFrame(val_encoded, columns=encoded_cols, index=None)
test_encoded_df = pd.DataFrame(test_encoded, columns=encoded_cols, index=None)

# Combine
X_train = pd.concat([X_train.drop(columns=nominal_cols), train_encoded_df], axis=1)
X_val = pd.concat([X_val.drop(columns=nominal_cols), val_encoded_df], axis=1)
X_test = pd.concat([X_test.drop(columns=nominal_cols), test_encoded_df], axis=1)

print(f"Final shapes - Train: {X_train.shape}, Val: {X_val.shape}")
```

Final shapes - Train: (1168, 200), Val: (292, 200)

```
In [38]: #checking for remaining NaN values
print("NaN counts after encoding:")
# Verify
print("Missing in X_train:", X_train.isnull().sum().sum())
print("Missing in X_val:", X_val.isnull().sum().sum())
print("Missing in X_test:", X_test.isnull().sum().sum())
```

NaN counts after encoding:

Missing in X_train: 14

Missing in X_val: 6

Missing in X_test: 10

```
In [39]: # Check which columns still have NaN
nan_cols_train = X_train.columns[X_train.isnull().any()].tolist()
nan_cols_val = X_val.columns[X_val.isnull().any()].tolist()
nan_cols_test = X_test.columns[X_test.isnull().any()].tolist()

print(f"NaN in X_train columns: {nan_cols_train}")
print(f"NaN in X_val columns: {nan_cols_val}")
print(f"NaN in X_test columns: {nan_cols_test}")
```

```
# Show sample of NaN values
if nan_cols_train:
    print("\nSample NaN rows in X_train:")
    print(X_train[X_train[nan_cols_train[0]].isnull()][nan_cols_train])

NaN in X_train columns: ['Functional']
NaN in X_val columns: ['Functional']
NaN in X_test columns: ['Functional']

Sample NaN rows in X_train:
    Functional
1013      NaN
710       NaN
542       NaN
666       NaN
1229      NaN
```

In [40]:

```
# Since it's ordinal, we fill NaN with 'Typ' (most common/most func
mode_val = X_train['Functional'].mode()[0]
print(f"\nMost common value: {mode_val}")

# Fill NaN
for df in [X_train, X_val, X_test]:
    df['Functional'] = df['Functional'].fillna(mode_val)

print("\nAfter filling:")
print(f"X_train NaN: {X_train['Functional'].isnull().sum()}")
print(f"X_val NaN: {X_val['Functional'].isnull().sum()}")
print(f"X_test NaN: {X_test['Functional'].isnull().sum()}")

# Verify ordinal encoding was applied correctly
print("\nFunctional unique values after filling:")
print(X_train['Functional'].unique())
```

Most common value: 8.0

After filling:
X_train NaN: 0
X_val NaN: 0
X_test NaN: 0

Functional unique values after filling:
[8. 6. 7. 5.]

Scale numerical features

In [41]:

```
# Identify numerical columns (non-binary, non-encoded)
numerical_cols = X_train.select_dtypes(include=['int64', 'float64'])

# Remove binary columns (already 0/1) and ordinal encoded columns
binary_and_ordinal = list(ordinal_mappings.keys()) + binary_cols
numerical_cols = [col for col in numerical_cols if col not in binary_and_ordinal]

print(f"Scaling {len(numerical_cols)} numerical columns")
```

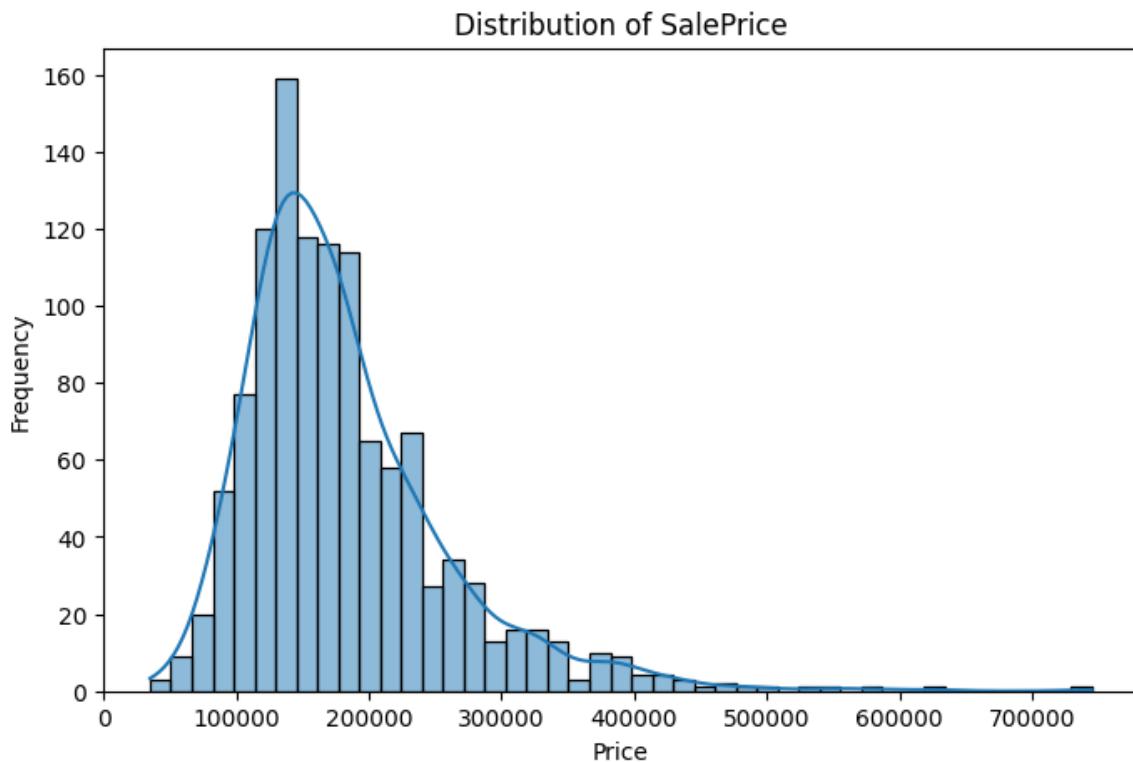
Scaling 180 numerical columns

```
In [42]: # Scale
scaler = StandardScaler()
X_train[numerical_cols] = scaler.fit_transform(X_train[numerical_cols])
X_val[numerical_cols] = scaler.transform(X_val[numerical_cols])
X_test[numerical_cols] = scaler.transform(X_test[numerical_cols])
```

Check target variable distribution (SalePrice)

The target variable, **SalePrice**, represents the final sale price of each house. We analyze its distribution to understand its range and shape. Since house prices typically exhibit right-skewed distributions, this step helps determine whether a transformation is required to improve model performance.

```
In [43]: # Histogram of the target
plt.figure(figsize=(8, 5))
sns.histplot(y_train, kde=True)
plt.title('Distribution of SalePrice')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()
```



```
In [44]: # Check skewness
print(f"Target skewness: {y_train.skew():.2f}")

# If skewed > 0.75, consider log transformation
if y_train.skew() > 0.75:
    y_train_log = np.log1p(y_train)
    y_val_log = np.log1p(y_val)
    print("Applied log transformation to target")
```

```
    else:
        y_train_log, y_val_log = y_train.copy(), y_val.copy()
```

Target skewness: 1.74
Applied log transformation to target

checking numeric feature distributions.

Highly skewed features can hurt linear models.we will use log transformation for skewed numeric features.

```
In [45]: # Check skewness of all numeric features
numeric_features = X_train.select_dtypes(include=[np.number]).columns
skewness = X_train[numeric_features].apply(lambda x: x.skew()).sort_values()

# Features with skewness > 0.75 (moderately skewed)
high_skew = skewness[abs(skewness) > 0.75]
print(f"Features with high skewness (>0.75): {len(high_skew)}")

# Apply log1p transformation to highly skewed features
for feature in high_skew.index:
    if X_train[feature].nunique() > 10:
        # Shift to positive values if needed
        min_val = X_train[feature].min()
        if min_val <= 0:
            shift = abs(min_val) + 1
            X_train[feature] = np.log1p(X_train[feature] + shift)
            X_val[feature] = np.log1p(X_val[feature] + shift)
            X_test[feature] = np.log1p(X_test[feature] + shift)
            print(f"Applied shifted log to: {feature} (min={min_val})")
        else:
            X_train[feature] = np.log1p(X_train[feature])
            X_val[feature] = np.log1p(X_val[feature])
            X_test[feature] = np.log1p(X_test[feature])
            print(f"Applied log to: {feature}")
```

Features with high skewness (>0.75): 168
Applied shifted log to: MiscVal (min=-0.09)
Applied shifted log to: LotArea (min=-0.87)
Applied shifted log to: 3SsnPorch (min=-0.12)
Applied shifted log to: LowQualFinSF (min=-0.12)
Applied shifted log to: BsmtFinSF2 (min=-0.29)
Applied shifted log to: ScreenPorch (min=-0.28)
Applied shifted log to: EnclosedPorch (min=-0.35)
Applied shifted log to: LotFrontage (min=-2.16)
Applied shifted log to: OpenPorchSF (min=-0.71)
Applied shifted log to: MasVnrArea (min=-0.60)
Applied shifted log to: BsmtFinSF1 (min=-0.97)
Applied shifted log to: TotalBsmtSF (min=-2.41)
Applied shifted log to: WoodDeckSF (min=-0.74)
Applied shifted log to: GrLivArea (min=-2.28)
Applied shifted log to: 1stFlrSF (min=-2.16)
Applied shifted log to: BsmtUnfSF (min=-1.28)
Applied shifted log to: 2ndFlrSF (min=-0.80)

Plotted distributions and skewness

We observe strong right skewness in features such as LotArea and GrLivArea, justifying the use of log transformation.

```
In [46]: # Keep only up to 20 features to plot
existing_features = [f for f in numeric_features if f in X_train.columns]
print(f'Plotting {len(existing_features)} features: {existing_features}')

fig, axes = plt.subplots(len(existing_features), 2, figsize=(12, 4*len(existing_features)))

for idx, feature in enumerate(existing_features):
    ax1, ax2 = axes[idx] if len(existing_features) > 1 else axes

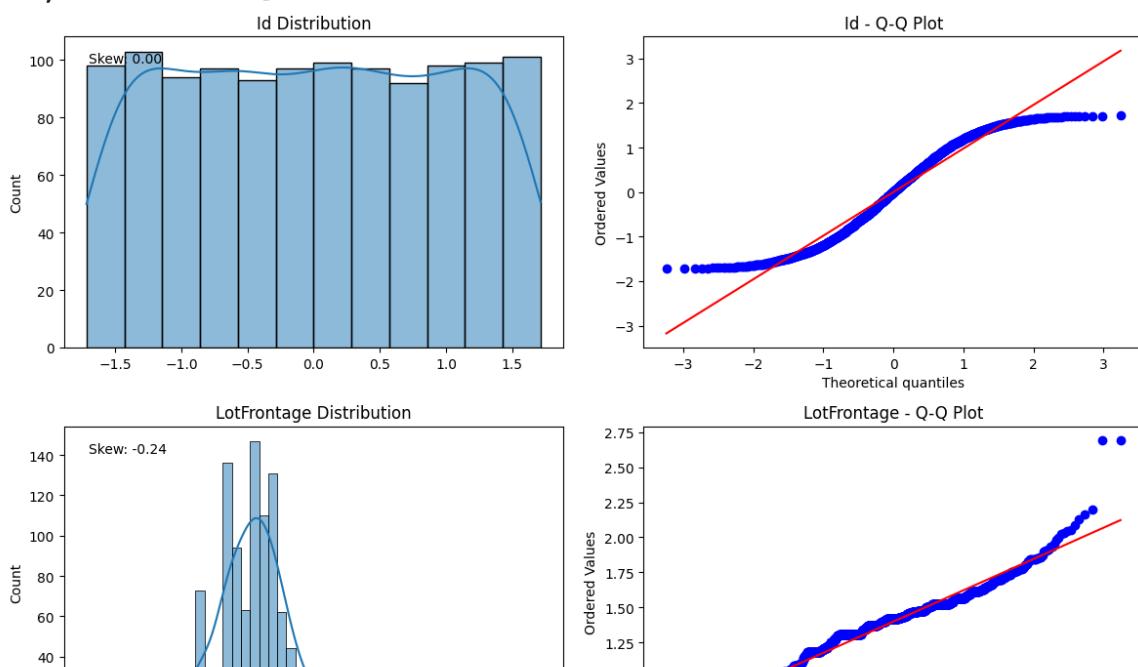
    # Histogram
    sns.histplot(X_train[feature], ax=ax1, kde=True)
    ax1.set_title(f'{feature} Distribution')
    ax1.set_xlabel('')

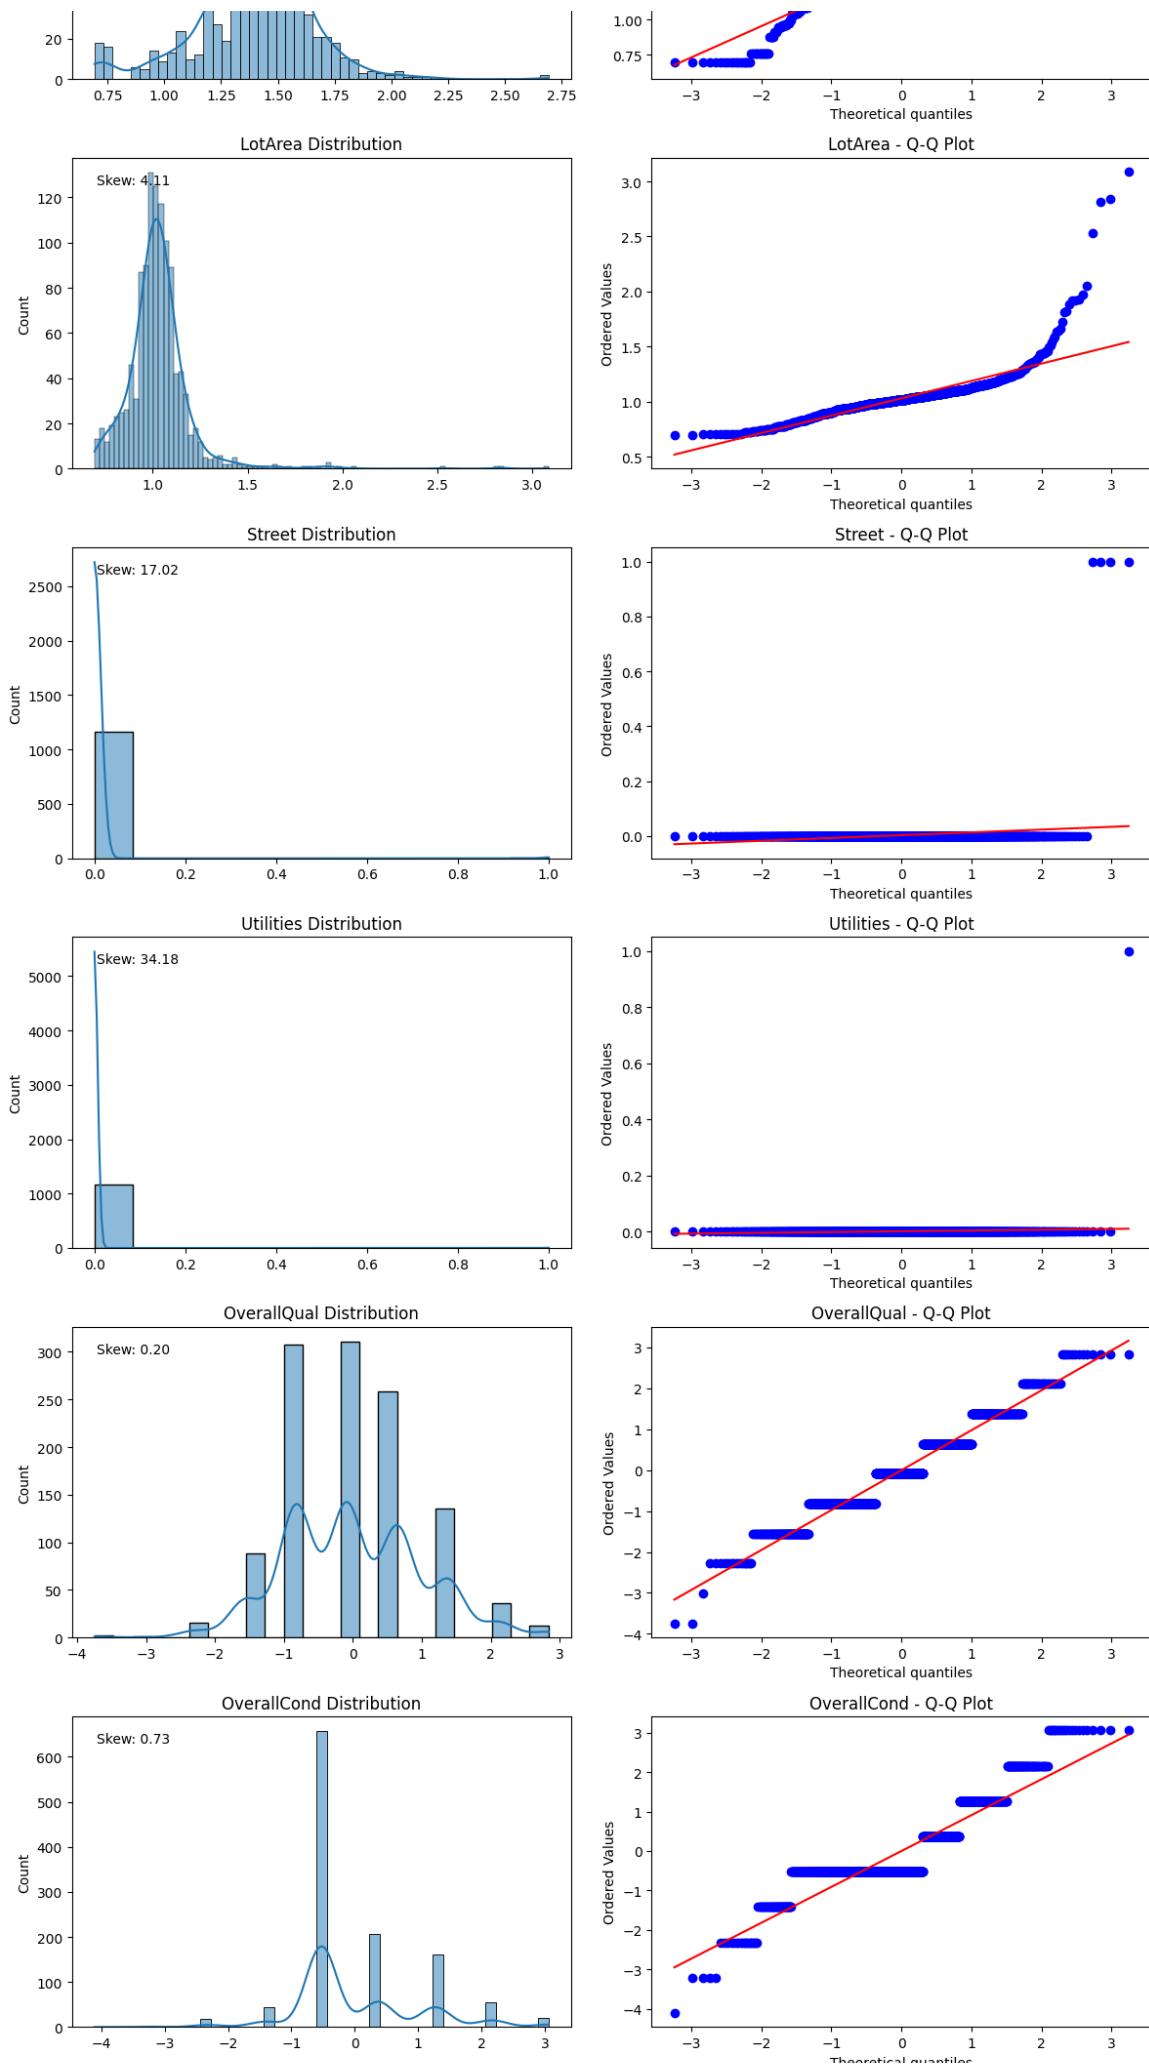
    # Q-Q plot
    from scipy import stats
    stats.probplot(X_train[feature].dropna(), dist="norm", plot=ax2)
    ax2.set_title(f'{feature} - Q-Q Plot')

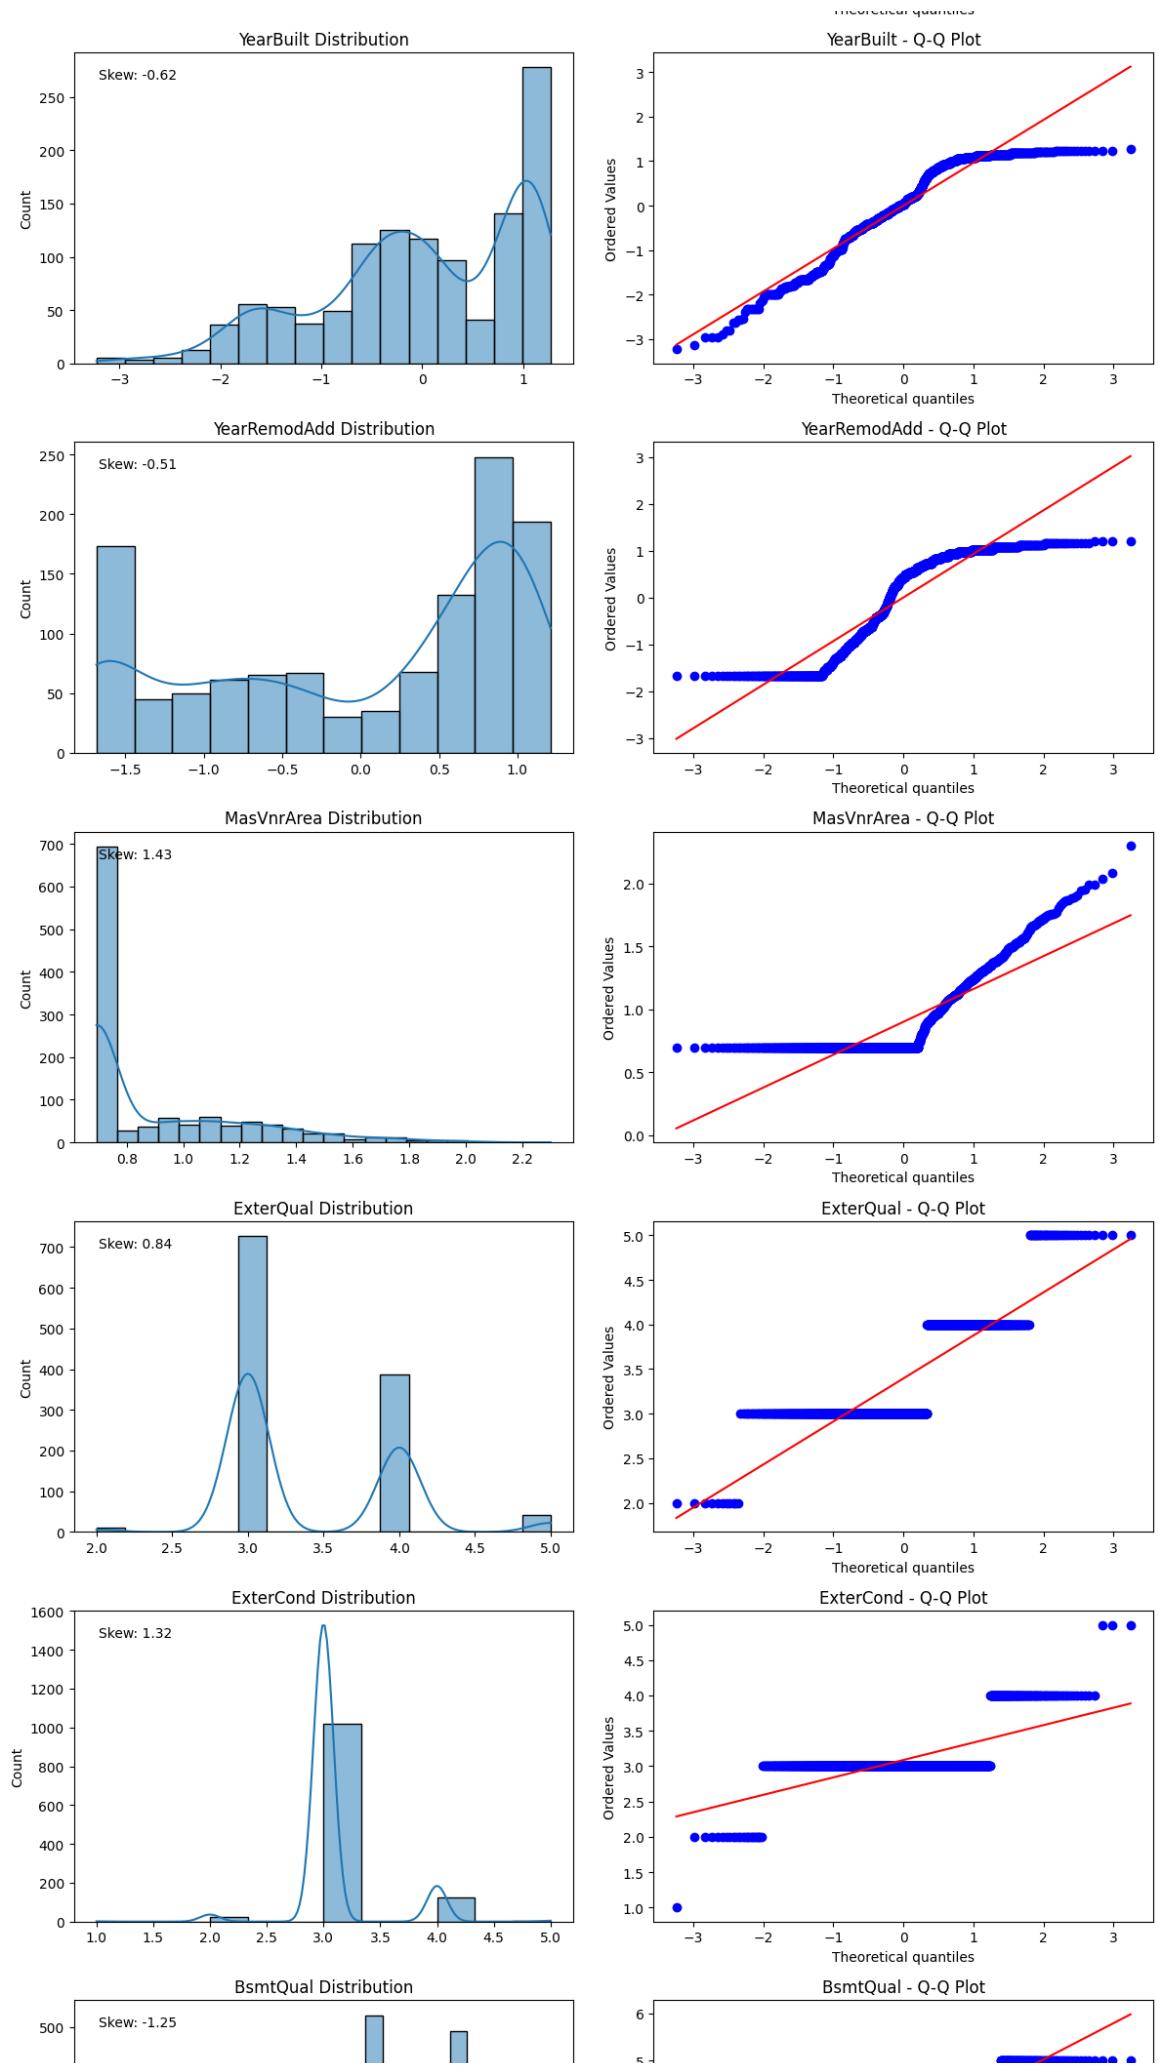
    # Show skewness
    skew_after = X_train[feature].skew()
    ax1.text(0.05, 0.95, f'Skew: {skew_after:.2f}', transform=ax1.transAxes, verticalalignment='top')

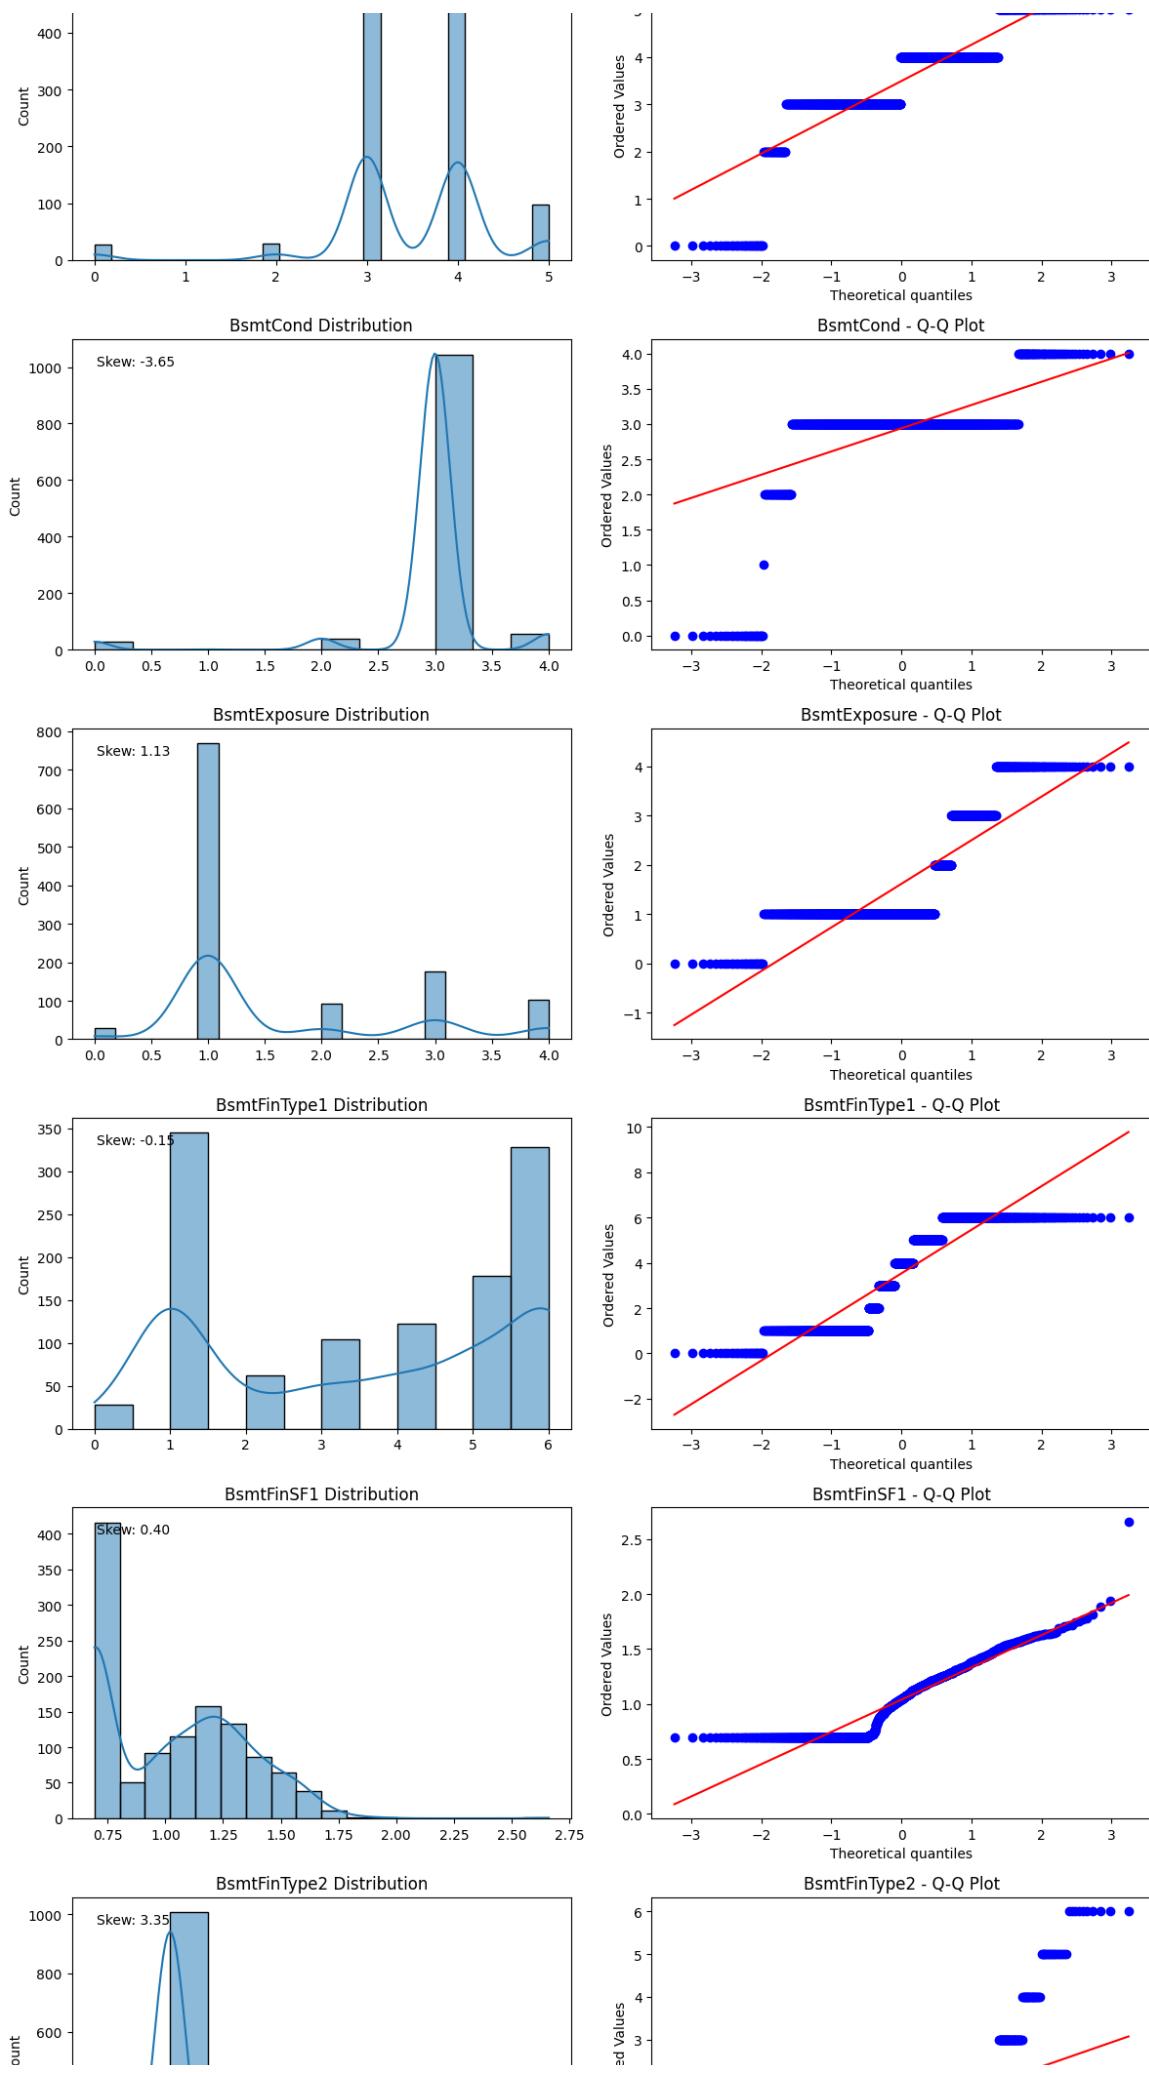
plt.tight_layout()
plt.show()
```

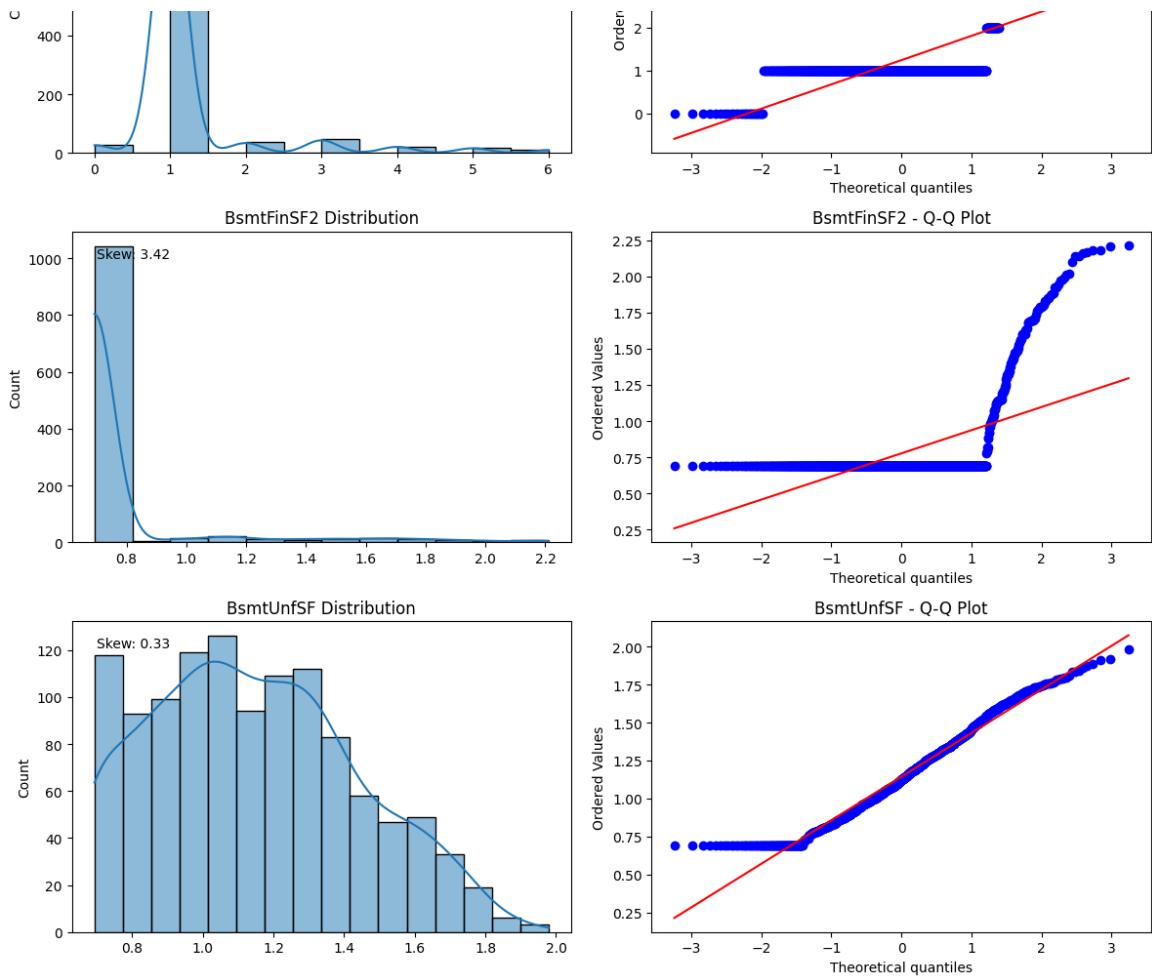
Plotting 20 features: ['Id', 'LotFrontage', 'LotArea', 'Street', 'Utilities', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'ExterQual', 'ExterCond', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF']











```
In [47]: import pickle
import os

save_path = "pkl/prepared_data.pkl"

with open(save_path, "wb") as f:
    pickle.dump({
        'X_train': X_train,
        'X_val': X_val,
        'X_test': X_test,
        'y_train': y_train,
        'y_val': y_val,
        'y_train_log': y_train_log, # Log transformed target
        'y_val_log': y_val_log,
        'encoder': encoder, # OneHotEncoder
        'scaler': scaler, # StandardScaler
        'nominal_cols': nominal_cols, # Column names for reference
        'numerical_cols': numerical_cols,
        'ordinal_mappings': ordinal_mappings,
        'binary_cols': binary_cols
    }, f)

print(f"Saved to {save_path}")
print(f"File exists: {os.path.exists(save_path)}")
print(f"File size: {os.path.getsize(save_path) / 1024 / 1024:.2f} M")
```

```
 Saved to pkl/prepared_data.pkl
```

```
 File exists: True
```

```
 File size: 4.56 MB
```

```
In [48]: # save_data = {  
#     'X_train': X_train,  
#     'X_val': X_val,  
#     'X_test': X_test,  
#     'y_train': y_train,  
#     'y_val': y_val,  
#     'y_train_log': y_train_log,  
#     'y_val_log': y_val_log,  
#     'all_columns': X_train.columns.tolist()  
# }
```

Conclusion

In this notebook, we performed comprehensive data preparation for the House Prices regression task. The dataset was explored, cleaned, and transformed using statistically sound techniques, including skewness.

By saving the prepared data, we ensure a reproducible and modular workflow that supports fair and consistent model comparison. This preprocessing stage establishes a strong foundation for training advanced regression models.

In the next notebook, **2-FeatureEngRegression.ipynb**, we will focus on Feature Engineering using the prepared dataset.